

# Reinforcement Learning for Restaurant Inventory & Preparation Management (Profit Objective)

## Research & Project Plan (Living Document)

Team

November 3, 2025

## Status & Scope

**Status:** Living document to be updated as design and experiments progress.

**Scope:** Build, evaluate, and document a reinforcement learning (RL) system that *jointly* orders raw ingredients and schedules in-house preparations to maximise expected discounted profit under budget, storage, lead-time, perishability (first-in-first-out; *FIFO*), and labour constraints. The policy must function across heterogeneous restaurants with variable menus and changing item sets.

## 1 Technical Implementation Sketch (from formulation notes)

### 1.1 Item taxonomy, demand coupling, and production

- **Two inventory types.** (i) *Raw ingredients* purchased from suppliers; (ii) *Preparations* (processed ingredients) produced in-house. Preparations consume raw inputs and *staff-hours* and have their own shelf-lives. We deplete stock via **FIFO** and track age-buckets.
- **Dish-driven demand.** We forecast per-dish demand  $\hat{d}_{k,t}$  and translate to inventory requirements using a bill-of-materials (**BOM**) matrix  $A \in \mathbb{R}^{N \times M}$ :

$$x_{i,t}^{\text{req}} = \sum_{k=1}^M A_{ik} \hat{d}_{k,t}.$$

- **Preparation production.** Daily prep decisions  $y_{j,t} \geq 0$  for preparation  $j$  consume raw inputs  $\sum_i B_{ij} y_{j,t}$  (prep-recipe matrix  $B$ ), use  $\tau_j$  staff-hours per unit, and yield new prep inventory at age 0.

### 1.2 MDP (*Markov decision process*) and state

Daily periodic review. State  $s_t$  concatenates global and item-wise components:

$$s_t = \left( \underbrace{\text{calendar}_t, B_t, C, \bar{H}_t}_{\text{global: budget, capacity, staff-hours}}, \underbrace{\{z_{i,t}\}_{i \in \mathcal{I}_t}}_{\text{variable set of item tokens}} \right),$$

where each item token  $z_{i,t}$  includes: on-hand age histogram, pipeline-by-days-to-arrival, purchase cost  $c_i$ , holding cost  $h_i$ , unit space  $u_i$ , case pack/minimum order multiple  $m_i$ , fixed order cost  $K_i$ , shelf-life  $L_i^{\text{life}}$ , price attribution  $p_i$  (from dish mix), recent sales, and forecast summaries

for implied usage. For preparations: labour  $\tau_i$  and yield parameters. The index set  $\mathcal{I}_t$  changes over time (items appear/disappear) and across restaurants.

**Permutation-invariant policy for an unfixed item set.** We use a DeepSets/attention encoder–decoder to handle a variable number of items:

$$D_i = \text{enc}_\theta(s_t^{\text{global}}, z_{i,t}), \quad g_t = \text{pool}(\{D_i\}), \quad \tilde{a}_{i,t} = \text{dec}_{\theta^I}(D_i, g_t),$$

with shared weights across items (*SKU*-wise parameter sharing; *SKU* = stock-keeping unit). The decoder emits continuous proposals  $(\tilde{q}_{i,t}, \tilde{y}_{i,t})$  for raw orders and prep batches.

**Alternative: multi-agent formulation.** As an alternative to the set-encoder, we can model each item  $i$  as an *agent* with parameter sharing across agents and a centralised critic (or value mixing). Communication/coordination can be implemented via attention over messages or mean-field signals. This naturally accommodates a variable item set, allows per-item exploration schedules, and enables explicit credit assignment, while keeping training stable through centralised training with decentralised execution.

### 1.3 Constraints and feasibility projection

Project proposals to feasible  $(q_{i,t}, y_{j,t})$ :

$$\sum_i c_i q_{i,t} \leq B_t, \quad \sum_i u_i(I_{i,t} + q_{i,t}) + \sum_j u_j^{\text{prep}} y_{j,t} \leq C, \quad \sum_j \tau_j y_{j,t} \leq \bar{H}_t, \quad q_{i,t} \in m_i \mathbb{Z}_+.$$

Projection uses a greedy knapsack routine ranking items by expected marginal profit-per-cost/per-space, with rounding to case packs. This stabilises training and guarantees constraint satisfaction.

### 1.4 Transitions with FIFO

Arrivals are dequeued from pipeline; new orders are enqueued at their lead-time slot. Preparation batches consume raw inputs via FIFO and create prep stock at age 0. Dish demand is sampled (Section 1.6); implied item usage is  $A D_t^{\text{dish}}$ ; sales are limited by available stock, consuming oldest ages first. Age advance and expiry produce spoilage write-offs.

### 1.5 Profit objective with explicit lost-sales penalty

Let  $S_{k,t}^{\text{dish}}$  be realised dish sales and  $\text{margin}_k$  the dish unit margin. Per-period reward:

$$\begin{aligned} r_t = & \underbrace{\sum_k p_k S_{k,t}^{\text{dish}}}_{\text{revenue}} - \underbrace{\sum_i c_i q_{i,t}}_{\text{purchasing}} + \underbrace{\sum_{i: q_{i,t} > 0} K_i}_{\text{holding}} - \underbrace{\sum_i h_i I_{i,t+1}}_{\text{holding}} - \underbrace{\sum_i \kappa_i \text{Spoil}_{i,t}}_{\text{spoilage}} \\ & - w \cdot \underbrace{\sum_j \tau_j y_{j,t}}_{\text{labour}} - \lambda \underbrace{\sum_k (d_{k,t} - S_{k,t}^{\text{dish}})_+}_{\text{lost-sales penalty}} \text{margin}_k, \end{aligned}$$

with  $\lambda \in [0, 1]$  (usually small since lost sales already reduce revenue). Set  $\kappa_i$  slightly above replacement cost to discourage waste.

## 1.6 Forecasting and simulator sampling

A light quantile forecaster (shared trunk; per-dish heads) outputs  $\{Q_{0.5}, Q_{0.8}, Q_{0.95}\}$  over a horizon  $H \geq \max_i L_i$ , trained with pinball loss. We use MC-dropout or small ensembles to obtain uncertainty. Simulator samples dish demand from these distributions conditional on calendar and events. Forecast summaries (means/quantiles/variance) are included in state so the policy is *forecast-aware*. This module is replaceable without touching the controller.

## 1.7 Algorithms

Actor-critic with: (i) global trunk for shared context; (ii) item towers with shared parameters; (iii) attention pooling to form a critic context. Start with PPO (proximal policy optimisation; stable on-policy) with entropy bonus; compare to SAC (soft actor-critic; off-policy continuous control). Optionally evaluate differentiable-simulator objectives (e.g., pathwise/DirectBack-prop) once the base PPO is stable.

## 2 Simulator: Pythonic Pseudocode (illustrative)

```

class KitchenEnv:
    def __init__(self, params, A_bom, B_prep, forecaster):
        self.A = A_bom          # [N x M] items <- dishes
        self.B = B_prep          # raw usage per prep unit
        self.fcst = forecaster  # dish-level quantile model with sampling
        self.cost = params.cost  # c_i, K_i, h_i, p_k, margin_k, kappa_i, w
        self.conf = params.conf   # L_i, m_i, u_i, shelf_i, capacity C
        self.hours = params.hours # tau_j, daily limit H_bar(t)
        self.budget = params.budget # B_t schedule
        self.calendar = params.calendar
        self._init_state()

    def _project(self, q_raw, y_prep):
        q_raw = round_to_casepacks(q_raw, m=self.conf.m)
        y_prep = cap_to_hours(y_prep, tau=self.hours.tau, H_bar=self.hours.today())
        q_raw, y_prep = knapsack_project(q_raw, y_prep,
                                         costs=self.cost, space=self.conf.u,
                                         C=self.conf.C, B=self.budget.today())
        return q_raw, y_prep

    def step(self, action):
        q, y = self._project(action.q, action.y)

        # Arrivals & pipeline rotation
        A = self.P[:, 0]
        self.P = shift_left(self.P)
        self.P[np.arange(self.N), self.conf.L-1] += q

        # Produce preparations (consume raws by FIFO; add prep at age 0)
        raw_need = self.B @ y
        self.I_age = fifo_consume(self.I_age, raw_need)  # raises if infeasible
        self.I_age[prep_ids, 0] += y

        # Sample dish demand and translate to item usage
        d_dish = self.fcst.sample(self.calendar.features())

```

```

use_items = self.A @ d_dish

# FIFO sales & ageing
S_items, I_next, spoil = fifo_sell_and_age(self.I_age, A, use_items,
                                             shelf=self.conf.shelf)
# Map S_items back to dish sales for revenue allocation if needed
S_dish = infer_dish_sales(S_items, self.A)

# Costs and reward
H_use = float(self.hours.tau @ y)
revenue = float(prices @ S_dish)
purchase = float(c @ q) + float((q > 0) @ K)
holding = float(h @ I_next.sum(axis=1))
lost_pen = float(margins @ np.maximum(d_dish - S_dish, 0.0)) * lambda_
spoilage = float(kappa @ spoil)

r = revenue - purchase - holding - spoilage - w * H_use - lost_pen
self.I_age = I_next
obs = self._make_obs()
done = self.calendar.is_terminal()
info = {"profit": r}
return obs, r, done, info

```

### 3 Baselines and Controls

1. **Heuristic (no optimisation): Demand + safety stock.** For each item, compute lead-time demand using forecast medians and add a fixed safety stock (e.g., based on historical service-level quantiles or a simple multiple of forecast  $\sigma$ ):

$$q_{i,t}^{\text{heur}} = \left[ (\hat{\mu}_{i,L} + z_\alpha \hat{\sigma}_{i,L}) - \text{stock\_position}_{i,t} \right]_+,$$

followed by budget/capacity-aware proportional downscaling and rounding to case packs. **This heuristic also serves as a *warm-start policy*** for RL by (a) initialising the actor via behaviour cloning on simulated trajectories, or (b) mixing its actions via  $\epsilon$ -greedy during early training.

2. **Static par levels (control).** Fixed days-of-cover targets per item (menu-category specific), adjusted seasonally; case-pack rounding only. No optimisation beyond feasibility checks.

### 4 Training Plan

#### Data

Daily dish sales; BOMs (dishes  $\rightarrow$  ingredients/preps); item costs  $c_i$ , dish prices  $p_k$  and margins; pack size  $m_i$ ; unit space  $u_i$ ; fixed order costs  $K_i$ ; shelf-lives; supplier lead-times; labour rates and prep labour  $\tau_j$ . Calendar features: day-of-week, week-of-year, holidays, events.

#### Curriculum and domain randomisation

Begin with small item sets, no fixed order costs, short lead-times, ample capacity. Gradually introduce  $K_i$ , tighter  $B_t, C, \bar{H}_t$ , longer lead-times, delivery uncertainty, and menu rotation (items

join/leave the set). Randomise costs/lead-times/budgets/capacity within 10 %–20 % to improve transfer across restaurants.

## Evaluation metrics

Episode profit, gross margin, fill-rate, stockout-days, inventory turns, waste rate, average and  $p95$  on-hand value, % capacity used, staff-hours, order frequency. Stress tests: holiday spikes; supplier delays.

## 5 Shadow/Offline Evaluation Protocol

Train forecaster on historical data; replay policy decisions offline using recorded receipts/sales; recompute KPIs subject to observed budgets/capacity/staff. Compare against baselines with paired tests and cost-component breakdowns.

## 6 Open Questions for Literature Review (to resolve with provided sources)

### 1. Set-encoder vs. multi-agent control.

- 1.a) Under what conditions (state/action coupling, scale, constraint structure) do permutation-invariant set encoders with shared towers outperform multi-agent (parameter-shared) policies with a central critic?
- 1.b) What coordination mechanisms (attention pooling, message passing, value mixing) are most sample-efficient for variable item sets and changing menus?
- 1.c) How does credit assignment compare (variance, stability) across the two formulations in inventory domains with hard feasibility projections?

### 2. Heuristic warm-start effectiveness.

- 2.a) Does behaviour cloning from a demand + safety-stock heuristic measurably improve early-stage learning curves, and what is the best mixture schedule (e.g.,  $\epsilon$ -greedy blending vs. supervised pretrain)?
- 2.b) Which safety-stock constructions (lead-time quantiles vs.  $(\mu + z_\alpha \sigma)$ ) provide the most effective initialisation under lost-sales and fixed order costs?

### 3. Profit shaping and lost-sales penalties.

- 3.a) What ranges of loss-scaling  $\lambda$  (relative to dish margin) lead to faster convergence without distorting the asymptotic profit optimum?
- 3.b) Are there principled potential-based shapings for profit (including spoilage and labour) that preserve optimal policies?

### 4. Perishability modelling granularity.

- 4.a) What age-bucket resolution is empirically sufficient for FIFO without incurring prohibitive state size?
- 4.b) Is waste cost better modelled as replacement cost, opportunity cost, or a calibrated overage factor to reflect operational externalities?

### 5. Feasibility projection design.

- 5.a) Which projection heuristics (greedy knapsack by margin-per-cost/per-space vs. proportional scaling) offer the best stability/return trade-off when embedded in the policy?
- 5.b) What rounding schemes to case packs minimise bias (e.g., stochastic rounding vs. deterministic nearest-multiple) under budget and capacity?

## 6. Preparation decisions and labour coupling.

- 6.a) How should prep batching and staff-hour costs be represented (continuous vs. integer batches) to balance realism and learnability?
- 6.b) Are there known policies for smoothing labour utilisation (e.g., variance penalties) that improve profit/waste without harming fill-rate?

## 7. Forecast-aware vs. forecast-free states.

- 7.a) Which forecast summaries (quantiles, moments, encoder tokens) most improve policy performance and robustness?
- 7.b) How sensitive is the policy to forecast misspecification, and which uncertainty treatments (MC-dropout, ensembles) mitigate degradation?

## 8. Differentiable simulation (*exo-MDP*) vs. standard RL.

- 8.a) In inventory tasks with exogenous demand and feasibility layers, when do pathwise/DirectBackprop objectives outperform PPO/SAC in sample efficiency and final return?
- 8.b) What reparameterisations are recommended for demand sampling and FIFO ageing to enable low-variance gradients?

## 9. Domain randomisation and cross-restaurant transfer.

- 9.a) What distributions and magnitudes of randomisation (lead times, MOQs, costs, capacity) yield the best out-of-domain generalisation?
- 9.b) How should restaurant/context embeddings be constructed (one-hot vs. learned features) for transfer without overfitting?

## 10. Evaluation protocol details.

- 10.a) Which metrics beyond profit (e.g., p95 on-hand value, waste rate, order frequency, labour variability) are most predictive of field success?
- 10.b) What stress scenarios (supplier delays, holiday spikes, menu rotations) are standardised in the literature for perishables?

## 11. Demand censoring and data issues.

- 11.a) How should censored sales (stockouts) be handled in forecaster training and offline evaluation to avoid optimistic/biased results?
- 11.b) Are there recommended filters or imputation strategies for building training windows with “healthy” on-hand conditions?

## Sources

#	Citation	ID (DOI / arXiv / name)	Summary (relevance)	Takeaways for our plan	Notes
1	Madeka, Torkkola, Eisenach, Luo, Foster, Kakade (2022), <i>Deep Inventory Management.</i>	arXiv:2210.03137	DRL with lost sales/lead-times; differentiable exogenous simulator and direct backprop show strong performance vs OR baselines.	Consider differentiable variants once PPO is stable; encode inventory structure in policy.	<a href="#">link</a>
2	Maggiar, Andaz, Bagaria, Eisenach, Foster, Gottesman, Perrault-Joncas (2025), <i>Structure-Informed DRL for Inventory.</i>	arXiv:2507.22040	Structure-aware policies (lead-times, stock position) and direct objectives across classical settings including perishables.	Use permutation-invariant item towers and structure tokens; aids generalisation.	<a href="#">link</a>
3	Sultana, Singh, et al. (2020), <i>Reinforcement Learning for Multi-Product Multi-Node Inventory Management.</i>	arXiv:2006.04037	Multi-product, capacity-coupled networks; actor-critic at scale.	Guides constraint handling and parameter sharing across large SKU sets.	<a href="#">link</a>
4	Kara & Doğan (2018), <i>Reinforcement learning approaches for determining the optimal ordering policy for perishable inventory systems. Expert Systems with Applications.</i>	10.1016/j.eswa.2017.08.046	Age-aware state and FIFO materially improve control for perishables; RL exceeds heuristics.	Include explicit age buckets and FIFO; model waste in reward.	<a href="#">link</a>
5	(EJOR Survey) <i>Reinforcement learning in inventory management: A roadmap</i> (2022).	10.1016/j.ejor.2021.07.016	Comprehensive survey; benchmarks, evaluation pitfalls, metrics.	Informs baselines, ablations, and robustness metrics.	<a href="#">link</a>
6	Kavoosi et al. (2025), <i>Dynamic pricing and perishable inventory management using deep RL. Expert Systems with Applications.</i>	Name: S095741742502189X	Joint pricing-ordering with perishables; continuous actions; vendor-managed context.	Blueprint for optional pricing extension with action pruning.	<a href="#">link</a>
7	Nomura & Liu (2025), <i>Deep RL for Dynamic Pricing and Perishable Inventory. Applied Sciences.</i>	10.3390/app15052421	Age-dependent demand with DRL; discusses invalid-action pruning.	Reinforces feasibility layer and action pruning ideas.	<a href="#">link</a>
8	Kaur & Prakash (2025), <i>Adaptive Inventory Strategies using DRL for Agri-Food.</i>	arXiv:2507.16670	DRL for perishable agri-food; uncertainty and robustness.	Stress-test patterns for perishables and uncertainty handling.	<a href="#">link</a>

#	Citation	ID (DOI / arXiv / name)	Summary (relevance)	Takeaways for our plan	Notes
9	(Tech/whitepaper) <i>End_to_End_revision3-10.pdf.</i>	Name: End_to_End_revision3-10	Industrial implementation details (pipelines/KPIs/offline testing).	Implementation templates for evaluation and deployment hygiene.	(no public link)
10	Medium (R. Ghosh), <i>Reinforcement Learning in Inventory Management: Leveraging AI for Optimal Order Management.</i>	Blog name	Pedagogical single-node sim and RL agents.	Simulator scaffolding and sanity-check baselines.	<a href="#">link</a>
11	Towards Data Science, <i>RL for Inventory Optimisation Series II: Multi-Echelon.</i>	Blog name	Multi-echelon example with PPO and action discretisation.	Reference patterns for networked settings and baselines.	<a href="#">link</a>
12	Medium (P. Kor), <i>Optimising Inventory Management with RL: A Hands-on Python Guide.</i>	Blog name	Q-learning tutorial; end-to-end code structure.	Quick prototype and unit-test fodder for the sim.	<a href="#">link</a>